



INSTITUTO SUPERIOR TÉCNICO

TRAFFIC ENGINEERING

METI

**Lab Report IV and V**

**Software Defined Networking and OpenFlow**

2018/2019

Group 6

André Mendes - 78079

Filipe Fernandes - 78083

# 1 Introduction

**TODO: Briefly describe both lab 4 and 5.**

This report is divided in two major sections. Laboratory 4 and 5.

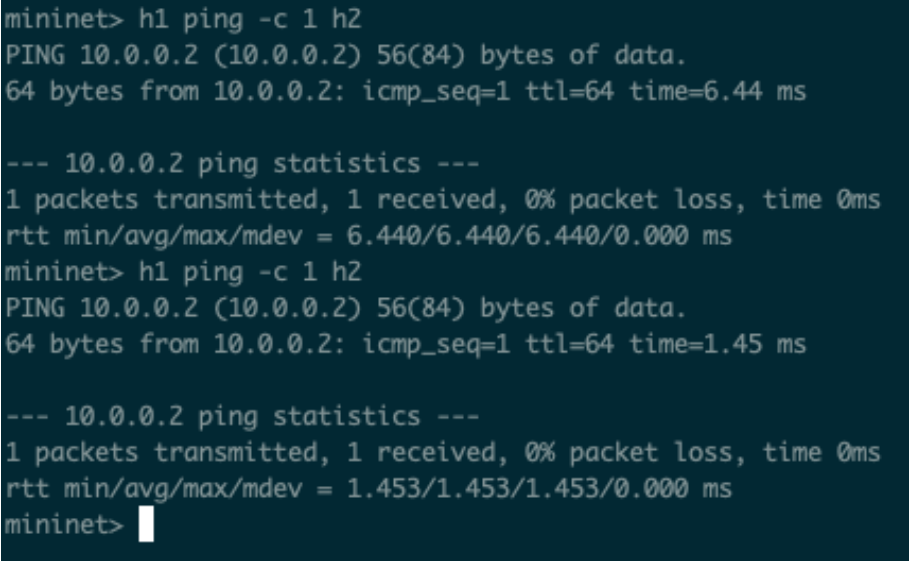
## 2 Laboratory 4

### 2.1 Mininet minimal topology

Mininet minimal topology is composed by 2 hosts, 1 switch and 1 basic controller. This topology can be launched with a single command from mininet like `sudo mn` or `sudo mn --topo=minimal`.

#### 2.1.1 Testing connectivity between hosts

On the next figure we can see the output of two consecutive pings from host 1 to host 2.



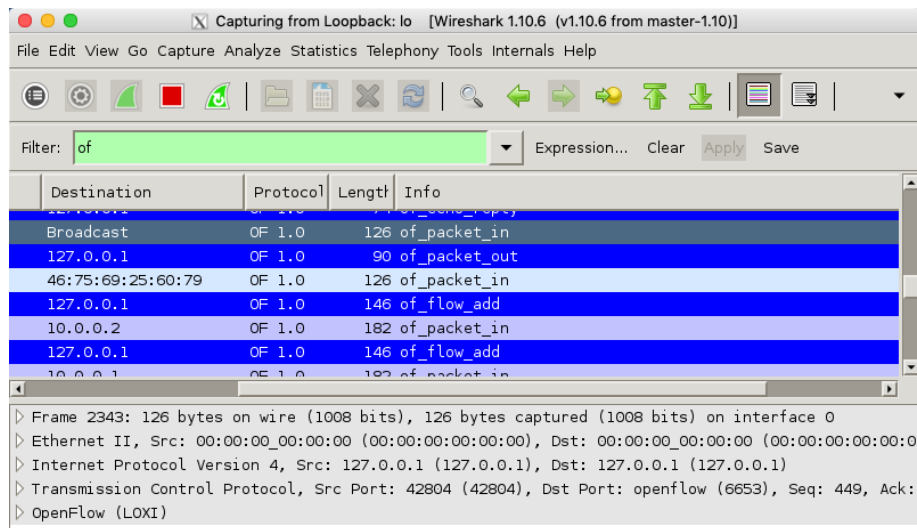
```
mininet> h1 ping -c 1 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=6.44 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 6.440/6.440/6.440/0.000 ms
mininet> h1 ping -c 1 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=1.45 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 1.453/1.453/1.453/0.000 ms
mininet> 
```

**Figure 1:** Ping from h1 to h2

As we can see the first one has a much higher delay than the second one. This happens because on the second ping a flow entry covering *ICMP* ping traffic was already installed in the switch, so no control traffic was generated, and the packets immediately pass through the switch. We can confirm this by looking at the packet trace captured during the first ping using *wireshark*, displayed on figure 2.

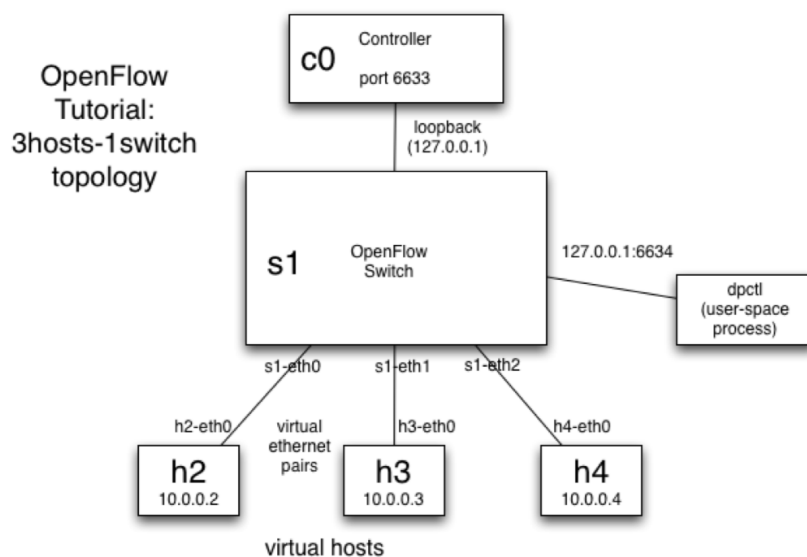


**Figure 2:** Wireshark capture on lo0 of mininet host during first ping

## 2.2 Mininet single topology

This topology includes 3 hosts, 1 OpenFlow Switch and 1 controller as displayed in figure 2.

### 2.2.1 Without controller



**Figure 3:** Mininet simple topology with 3 hosts, hosts actually are h1(10.0.0.1), h2(10.0.0.2) and h3(10.0.0.3) instead of h2, h3 and h4

Right after deploying this topology the switch flow table is empty. So if we try to ping it will fail because the switch hasn't any configurations and it doesn't even act as a hub.

If we add the following flows on the switch flow-table it will start doing what we expected firstly.

```

sudo ovs-ofctl add-flow s1 in_port=1,actions=output:2
sudo ovs-ofctl add-flow s1 in_port=2,actions=output:1

```

We are configuring the switch to redirect traffic from input port 1 to output port 2 and vice-versa. The content of the flow-table after these commands are displayed on the next figure.

```
mininet@mininet-vm:~$ sudo ovs-ofctl add-flow s1 in_port=1,actions=output:2
mininet@mininet-vm:~$ sudo ovs-ofctl add-flow s1 in_port=2,actions=output:1
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
 cookie=0x0, duration=12.447s, table=0, n_packets=0, n_bytes=0, idle_age=12, in_port=1 actions=output:2
 cookie=0x0, duration=6.464s, table=0, n_packets=0, n_bytes=0, idle_age=6, in_port=2 actions=output:1
mininet@mininet-vm:~$
```

**Figure 4:** Install flows on switch s1

And now we can successfully ping between h1 and h2 as seen on the next figure.

```
mininet> h1 ping -c3 h2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.
 64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.225 ms
 64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.054 ms
 64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.052 ms

--- 10.0.0.2 ping statistics ---
 3 packets transmitted, 3 received, 0% packet loss, time 2002ms
 rtt min/avg/max/mdev = 0.052/0.110/0.225/0.081 ms
mininet>
```

**Figure 5:** Ping between h1 and h2 with flows installed

After each traffic flow that belongs to the flows installed on s1, the flow-table of s1 is updated.

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
 cookie=0x0, duration=13.665s, table=0, n_packets=4, n_bytes=336, idle_age=2, in_port=1 actions=output:2
 cookie=0x0, duration=11.026s, table=0, n_packets=4, n_bytes=336, idle_age=2, in_port=2 actions=output:1
mininet@mininet-vm:~$
```

**Figure 6:** Flow table of s1 after ping

All pings executed between other hosts than h1 to h2 and h2 to h1 will fail. Except if we add/edit the flows for other routes. If we delete the current flows the connection between h1 and h2 will also start to fail.

### 2.2.2 With controller

If we repeat the steps above but this time with a controller, the behaviour of the network is different.

Let's start by launching a controller on mininet: `$ controller ptcp:6653`.

When the controller starts a lot of messages are exchanged.

These messages are:

- Hello - (Controller<->Switch)
- Features - (RequestController->Switch)
- Set Config - (RequestController->Switch)
- Features Reply - (Switch->Controller)
- Port Status - (Switch->Controller)

Some of them are captured on figure 7.

No.	Time	Source	Destination	Protocol	Length	Info
33315	3546.033797	127.0.0.1	127.0.0.1	OF 1.0	74	of_hello
33317	3546.034038	127.0.0.1	127.0.0.1	OF 1.0	74	of_hello
33319	3546.034299	127.0.0.1	127.0.0.1	OF 1.0	74	of_features_req
33321	3546.034368	127.0.0.1	127.0.0.1	OF 1.0	290	of_features_repl
33322	3546.034452	127.0.0.1	127.0.0.1	OF 1.0	78	of_set_config
33354	3551.034143	127.0.0.1	127.0.0.1	OF 1.0	74	of_echo_request
33355	3551.034344	127.0.0.1	127.0.0.1	OF 1.0	74	of_echo_reply
33387	3556.033416	127.0.0.1	127.0.0.1	OF 1.0	74	of_echo_request
33388	3556.033620	127.0.0.1	127.0.0.1	OF 1.0	74	of_echo_reply
33426	3560.034043	127.0.0.1	127.0.0.1	OF 1.0	74	of_hello

**Figure 7:** Wireshark capture of messages being exchanged between controller and switch.

When we generate packets (h1 ping -c1 h2) new messages are exchanged. This messages can be:

- Packet-In (Switch->Controller)
- Packet-Out (Controller->Switch)
- Flow-Mod (Controller->Switch)
- Flow-Expired (Switch->Controller)

On figure 8 some of these are captured. This is an example of OpenFlow in reactive mode (flows are pushed down in response to individual packets). The other possible mode is proactive mode where flows can be pushed down before packets to avoid the round-trip times and flow insertion delays.

No.	Time	Source	Destination	Protocol	Length	Info
46190	3790.308888	00:00:00_00:00:01	Broadcast	OF 1.0	126	of_packet_in
46191	3790.309255	127.0.0.1	127.0.0.1	OF 1.0	90	of_packet_out
46193	3790.309546	00:00:00_00:00:02	00:00:00_00:00:01	OF 1.0	126	of_packet_in
46194	3790.309653	127.0.0.1	127.0.0.1	OF 1.0	146	of_flow_add
46195	3790.310122	10.0.0.1	10.0.0.2	OF 1.0	182	of_packet_in
46196	3790.310267	127.0.0.1	127.0.0.1	OF 1.0	146	of_flow_add
46197	3790.310729	10.0.0.2	10.0.0.1	OF 1.0	182	of_packet_in
46198	3790.310835	127.0.0.1	127.0.0.1	OF 1.0	146	of_flow_add
46417	3795.312858	00:00:00_00:00:02	00:00:00_00:00:01	OF 1.0	126	of_packet_in
46418	3795.313135	127.0.0.1	127.0.0.1	OF 1.0	146	of_flow_add
46420	3795.313474	00:00:00_00:00:01	00:00:00_00:00:02	OF 1.0	126	of_packet_in
46421	3795.313565	127.0.0.1	127.0.0.1	OF 1.0	146	of_flow_add

**Figure 8:** Wireshark capture of flows being pushed down in response to individual packets.

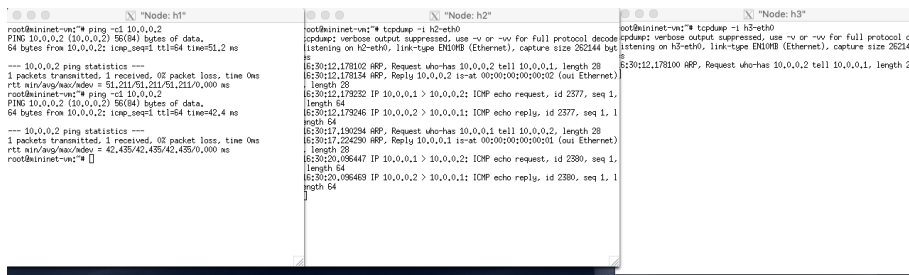
**TODO: Explain this new messages**

### 2.2.3 Benchmark controller with iperf

To benchmark the reference controller we will use iperf, a tool used to check speeds between two computers.

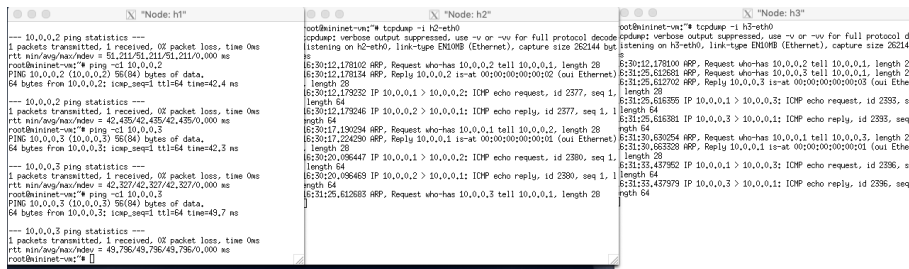
On the next figure we can see the results of iperf between all hosts. The higher speed obtained is 24.7 Gbits/s on link h1<->h2. But all hosts have about the same speed between them, around 24Gbits/s.





**Figure 11:** Output of ping in h1 to h2 and tcpdump on h2 and h3

A similar behaviour is verified when we perform another two pings from h1 to h3, only the first is flooded to h2. As seen on figure 12.




**Figure 12:** Output of ping in h1 to h2, h3 and tcpdump on h2 and h3

On the following image we present a wireshark print, captured on mininet VM, interface lo0 that confirms all packets are sent to controller.

1741	30.246244000	00:00:00_00:00:01	Broadcast	OF 1.0	126 of_packet_in
1742	30.262423000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
1744	30.262787000	00:00:00_00:00:02	00:00:00_00:00:01	OF 1.0	126 of_packet_in
1745	30.263304000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
1746	30.263491000	10.0.0.1	10.0.0.2	OF 1.0	182 of_packet_in
1747	30.263844000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
1748	30.264028000	10.0.0.2	10.0.0.1	OF 1.0	182 of_packet_in
1749	30.264525000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
2312	34.430468000	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
2313	34.430856000	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
2341	35.277381000	00:00:00_00:00:02	00:00:00_00:00:01	OF 1.0	126 of_packet_in
2342	35.304876000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
2344	35.305166000	00:00:00_00:00:01	00:00:00_00:00:02	OF 1.0	126 of_packet_in
2345	35.305420000	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out

**Figure 13:** Wireshark capture on interface Loopback0 with "of" filter on mininet VM

And also as expected no routes are installed at s1, but mac to port table is properly learned. See figures 14 and 15



**Figure 14:** Switch 1 flow table entries

```

mininet-vm: (ssh) M1 ~ _inet-vm:~(pox) ssh M2 ~ _mininet-vm: (ssh) M3 ~ _T/Labs/Lab4.5 (ssh) M4
DEBUG:misc.of.switch.with_flow:Installing flow...
DEBUG:misc.of.switch.with_flow:Source: 00:00:00:00:00:02 Destination: 00:00:00:00:00:03 Port: 3
ACINFO:core:Going down...
INFO:openflow.of_01:[00:00:00:00:00:01 1] disconnected
INFO:core:Down...
mininetmininet-vm:~(pox) $ ./pox.py log.level --DEBUG misc.of.switch.without_flow
POX 0.2.0 (carp) / Copyright 2011-2013 James McCauley, et al.
DEBUG:core:POX 0.2.0 (carp) going up...
DEBUG:core:Running on Cython (2.7.6 Oct. 26 2016 20:30:19)
DEBUG:core:Platform is Linux-4.2.0-27-generic-x86_64-with-Ubuntu-14.04-trusty
INFO:core:POX 0.2.0 (carp) is up.
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633
INFO:openflow.of_01:[00:00:00:00:00:01 1] connected
DEBUG:misc.of.switch.without_flow:Controlling [00:00:00:00:00:01 1]
ACINFO:core:Going down...
INFO:openflow.of_01:[00:00:00:00:00:01 1] disconnected
INFO:core:Down...
mininetmininet-vm:~(pox) $ ./pox.py log.level --DEBUG misc.of.switch.without_flow
POX 0.2.0 (carp) / Copyright 2011-2013 James McCauley, et al.
DEBUG:core:POX 0.2.0 (carp) going up...
DEBUG:core:Running on Cython (2.7.6 Oct. 26 2016 20:30:19)
DEBUG:core:Platform is Linux-4.2.0-27-generic-x86_64-with-Ubuntu-14.04-trusty
INFO:core:POX 0.2.0 (carp) is up.
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633
INFO:openflow.of_01:[00:00:00:00:00:01 1] connected
DEBUG:misc.of.switch.without_flow:Controlling [00:00:00:00:00:01 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]
DEBUG:misc.of.switch.without_flow:[EthAddr('00:00:00:00:00:02'): 2, EthAddr('00:00:00:00:00:03'): 3, EthAddr('00:00:00:00:00:01'): 1]

```

Figure 15: Log of POX with all mac to port table entries being displayed

### 3.3 Implementing a learning switch with table flow entries

On this subsection, the behaviour is similar to the previous but now table flow entries are installed at s1. This means that most of the packets won't go through the controller. As we will confirm.

We started by performing the same pings as we did before, two consecutive pings from h1 to h2. As expected only the first is flooded to h3 and sent to the controller. See figures 16 and 17.

```

root@mininet-vm:~# ping -c 1 10.0.0.2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.431 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0.0 ms
rtt min/avg/mdev = 0.431/0.431/0.431 ms

root@mininet-vm:~# ping -c 1 10.0.0.2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.431 ms

--- 10.0.0.2 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0.0 ms
rtt min/avg/mdev = 0.431/0.431/0.431 ms

root@mininet-vm:~#

root@mininet-vm:~# tcpdump -i h2-eth0
tcpdump: verbose output suppressed, use -v or -vv for full protocol decoding
listening on h2-eth0, link-type EN10MB (Ethernet), capture size 262144
68.3545.242409 ARP, Request who-has 10.0.0.2 tell 10.0.0.1, length 28
68.3545.242443 ARP, Reply 10.0.0.2 is-at 00:00:00:00:00:02 (oui Ethernet)
length 28
68.3545.242639 IP 10.0.0.1 > 10.0.0.2: ICMP echo request, id 2477, seq 1, length 64
68.3545.242656 IP 10.0.0.2 > 10.0.0.1: ICMP echo reply, id 2477, seq 1, length 64
68.3545.244165 ARP, Request who-has 10.0.0.1 tell 10.0.0.2, length 28
68.3545.244612 ARP, Reply 10.0.0.1 is-at 00:00:00:00:00:01 (oui Ethernet)
length 28
68.3545.137117 IP 10.0.0.1 > 10.0.0.2: ICMP echo request, id 2481, seq 1, length 64
68.3545.137132 IP 10.0.0.2 > 10.0.0.1: ICMP echo reply, id 2481, seq 1, length 64

```

Figure 16: Output of ping in h1 to h2 and tcpdump on h2 and h3

219	11.50071500X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
220	11.50921800X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
251	14.84899800X	00:00:00_00:00:01	Broadcast	OF 1.0	126 of_packet_in
252	14.88860100X	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
254	14.88923000X	00:00:00_00:00:02	00:00:00_00:00:01	OF 1.0	126 of_packet_in
256	14.89006000X	127.0.0.1	127.0.0.1	OF 1.0	90 of_packet_out
257	14.89029800X	10.0.0.1	10.0.0.2	OF 1.0	182 of_packet_in
258	14.89516200X	127.0.0.1	127.0.0.1	OF 1.0	146 of_flow_add
273	14.93252100X	127.0.0.1	127.0.0.1	OF 1.0 +	170 of_packet_out + of_flow_add
335	19.50192400X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
336	19.52476800X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
382	24.50273500X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
384	24.55182000X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
482	29.50177500X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
484	29.52901600X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
525	34.50205000X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
526	34.50853300X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
561	39.50203500X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
562	39.53653200X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
597	44.50170600X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
598	44.51625800X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply
633	49.50158800X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_request
635	49.54198500X	127.0.0.1	127.0.0.1	OF 1.0	74 of_echo_reply

Figure 17: Wireshark capture on interface Loopback0 with "of" filter on mininet VM, displaying the addition of a flow and packet re-route



This happens because the first time packets are sent to the controller, flow table entries are installed at s1. See figures 18 and 19.

```
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633
INFO:openflow.of_01:[00-00-00-00-00-01 1] connected
DEBUG:misc.of_switch_with_flow:Controlling [00-00-00-00-00-01 1]
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:02 Destination: 00:00:00:00:00:01 Port: 1
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:01 Destination: 00:00:00:00:00:02 Port: 2
```

**Figure 18:** Log of running POX displaying messages when installing flows on s1

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
 cookie=0x0, duration=36.698s, table=0, n_packets=3, n_bytes=238, idle_age=26, in_port=2,dl_dst=00:00:00:00:00:01 actions=output:1
 cookie=0x0, duration=36.66s, table=0, n_packets=2, n_bytes=140, idle_age=26, in_port=1,dl_dst=00:00:00:00:00:02 actions=output:2
mininet@mininet-vm:~$
```

**Figure 19:** Flows installed at s1 after ping from h1 to h2

On the next figure we can also confirm that mac to port table is properly learned.

```
mininet@mininet-vm:~/pox$ ./pox.py log.level --DEBUG misc.of_switch_with_flow
POX 0.2.0 (carp) / Copyright 2011-2013 James McCauley, et al.
DEBUG:core:POX 0.2.0 (carp) going up...
DEBUG:core:Running on (Python (2.7.6/Oct 26 2016 20:30:19))
DEBUG:core:Platform is Linux-4.2.0-27-generic-x86_64-with-Ubuntu-14.04-trusty
INFO:core:POX 0.2.0 (carp) is up.
DEBUG:openflow.of_01:Listening on 0.0.0.0:6633
INFO:openflow.of_01:[00-00-00-00-00-01 1] connected
DEBUG:misc.of_switch_with_flow:Controlling [00-00-00-00-00-01 1]
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:02')}: 2, EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:02 Destination: 00:00:00:00:00:01 Port: 1
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:02')}: 2, EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:01 Destination: 00:00:00:00:00:02 Port: 2
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:02')}: 2, EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:02')}: 2, EthAddr('00:00:00:00:00:03')}: 3, EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:03 Destination: 00:00:00:00:00:01 Port: 1
DEBUG:misc.of_switch_with_flow:{EthAddr('00:00:00:00:00:02')}: 2, EthAddr('00:00:00:00:00:03')}: 3, EthAddr('00:00:00:00:00:01')}: 1}
DEBUG:misc.of_switch_with_flow:Installing flow...
DEBUG:misc.of_switch_with_flow:Source: 00:00:00:00:00:01 Destination: 00:00:00:00:00:03 Port: 3
```

**Figure 20:** Log of POX with all mac to port table entries being displayed

On figure 21 we present the final table flow entry of s1 after performing pings from all to all. Our flows have two matching conditions and one output action.

- `msg.match.in_port = packet.in.in_port`
- `msg.match.dl_dst = packet.dst`
- `msg.actions.append(of.ofp_action_output(port = self.mac_to_port.get(packet.dst)))`

Summarizing, packets from a certain port to a certain destination are re-routed to the respective destination port.

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
 cookie=0x0, duration=31.314s, table=0, n_packets=3, n_bytes=238, idle_age=21, in_port=3,dl_dst=00:00:00:00:00:02 actions=output:2
 cookie=0x0, duration=1749.538s, table=0, n_packets=7, n_bytes=518, idle_age=32, in_port=2,dl_dst=00:00:00:00:00:01 actions=output:1
 cookie=0x0, duration=1747.159s, table=0, n_packets=6, n_bytes=420, idle_age=18, in_port=3,dl_dst=00:00:00:00:00:01 actions=output:1
 cookie=0x0, duration=31.277s, table=0, n_packets=2, n_bytes=140, idle_age=21, in_port=2,dl_dst=00:00:00:00:00:03 actions=output:3
 cookie=0x0, duration=1747.121s, table=0, n_packets=5, n_bytes=322, idle_age=18, in_port=1,dl_dst=00:00:00:00:00:03 actions=output:3
 cookie=0x0, duration=1749.501s, table=0, n_packets=6, n_bytes=420, idle_age=32, in_port=1,dl_dst=00:00:00:00:00:02 actions=output:2
mininet@mininet-vm:~$
```

**Figure 21:** Flows installed at s1 after ping from all to all

## 4 Conclusions

Regarding laboratory 4 our goal was mainly to get hands on Mininet and OpenFlow protocol. We found mininet very interesting because it allows us to use kernel linux/unix functions with just a few commands.

We now understand main benefits of SDN (Software Defined Networks) and how it can help us create an abstraction layer between the equipment's configuration and the topology of the network.

On laboratory 5 we had the opportunity to test POX, a python framework for communicating with SDN switches using OpenFlow. We started by analysing the behaviour of an Hub, followed by a learning switch with and without table flow entries. In all cases we performed pings to test connectivity and RTT(round trip times). We conclude that the lower RTT happens when using learning switch with table flow entries (after the flows being inserted on switch, values around 0.431ms). And the higher is when using a learning switch with a controller but without installing table flow entries (about 49 ms). The Hub behaviour stays in the middle when comparing RTTs ( +/- 7 ms).

## Annex:

### A

#### Annex A - of\_switch\_without\_flow.py

```
# Copyright 2012 James McCauley
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at:
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
    implied.
# See the License for the specific language governing permissions and
# limitations under the License.
```

```
"""
```

```
This component is for use with the OpenFlow tutorial.
```

```
It acts as a simple hub, but can be modified to act like an L2
learning switch.
```

```
It's roughly similar to the one Brandon Heller did for NOX.
"""
```

```
from pox.core import core
import pox.openflow.libopenflow_01 as of
```

```
log = core.getLogger()
```

```
class Tutorial (object):
```

```
"""
```

```
A Tutorial object is created for each switch that connects.
A Connection object for that switch is passed to the __init__
function.
"""
```

```
def __init__ (self, connection):
```

```
# Keep track of the connection to the switch so that we can
# send it messages!
    self.connection = connection
```

```
# This binds our PacketIn event listener
    connection.addListener(self)
```

```
# Use this table to keep track of which ethernet address is on
```

```

# which switch port (keys are MACs, values are ports).
self.mac_to_port = {}

def resend_packet (self, packet_in, out_port):
    """
    Instructs the switch to resend a packet that it had sent to us.
    "packet_in" is the ofp_packet_in object the switch had sent to
    the
    controller due to a table-miss.
    """
    msg = of.ofp_packet_out()
    msg.data = packet_in

    # Add an action to send to the specified port
    action = of.ofp_action_output(port = out_port)
    msg.actions.append(action)

    # Send message to switch
    self.connection.send(msg)

def act_like_hub (self, packet, packet_in):
    """
    Implement hub-like behavior — send all packets to all ports
    besides
    the input port.
    """

    # We want to output to all ports — we do that using the special
    # OFPP_ALL port as the output port. (We could have also used
    # OFPP_FLOOD.)
    self.resend_packet(packet_in, of.OFPP_ALL)

    # Note that if we didn't get a valid buffer_id, a slightly better
    # implementation would check that we got the full data before
    # sending it (len(packet_in.data) should be == packet_in.
    # total_len)).

def act_like_switch (self, packet, packet_in):
    """
    Implement switch-like behavior.
    """

    # Learn the port for the source MAC
    self.mac_to_port.update({ packet.src: packet_in.in_port })

    log.debug(self.mac_to_port)

    if self.mac_to_port.get(packet.dst) is not None:
        # Send packet out the associated port

```

```

        self.resend_packet(packet_in, self.mac_to_port.get(packet.dst
        ))
    else:
        self.resend_packet(packet_in, of.OFPP_ALL)

def _handle_PacketIn (self, event):
    """
    Handles packet in messages from the switch.
    """

    packet = event.parsed # This is the parsed packet data.
    if not packet.parsed:
        log.warning("Ignoring_incomplete_packet")
        return

    packet_in = event.ofp # The actual ofp_packet_in message.

    self.act_like_switch(packet, packet_in)

def launch ():
    """
    Starts the component
    """
    def start_switch (event):
        log.debug("Controlling_%s" % (event.connection,))
        Tutorial(event.connection)
    core.openflow.addListenerByName("ConnectionUp", start_switch)

```

## B

### Annex B - of\_switch\_with\_flow.py

```

# Copyright 2012 James McCauley
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at:
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
# implied.
# See the License for the specific language governing permissions and
# limitations under the License.
"""
This component is for use with the OpenFlow tutorial.

```

*It acts as a simple hub, but can be modified to act like an L2 learning switch.*

*It's roughly similar to the one Brandon Heller did for NOX.*  
"""

```
from pox.core import core
import pox.openflow.libopenflow_01 as of
```

```
log = core.getLogger()
```

```
class Tutorial (object):
```

```
    """
```

```
    A Tutorial object is created for each switch that connects.
    A Connection object for that switch is passed to the __init__
    function.
    """
```

```
    def __init__ (self, connection):
```

```
        # Keep track of the connection to the switch so that we can
        # send it messages!
```

```
        self.connection = connection
```

```
        # This binds our PacketIn event listener
```

```
        connection.addListener(self)
```

```
        # Use this table to keep track of which ethernet address is on
        # which switch port (keys are MACs, values are ports).
```

```
        self.mac_to_port = {}
```

```
    def resend_packet (self, packet_in, out_port):
```

```
        """
```

```
        Instructs the switch to resend a packet that it had sent to us.
        "packet_in" is the ofp_packet_in object the switch had sent to
        the
        controller due to a table-miss.
        """
```

```
        msg = of.ofp_packet_out()
```

```
        msg.data = packet_in
```

```
        # Add an action to send to the specified port
```

```
        action = of.ofp_action_output(port = out_port)
```

```
        msg.actions.append(action)
```

```
        # Send message to switch
```

```
        self.connection.send(msg)
```

```
    def act_like_hub (self, packet, packet_in):
```

```
        """
```

```

Implement hub-like behavior — send all packets to all ports
besides
the input port.
"""

# We want to output to all ports — we do that using the special
# OFPP_ALL port as the output port. (We could have also used
# OFPP_FLOOD.)
self.resend_packet(packet_in, of.OFPP_ALL)

# Note that if we didn't get a valid buffer_id, a slightly better
# implementation would check that we got the full data before
# sending it (len(packet_in.data) should be == packet_in.
total_len)).

def act_like_switch (self, packet, packet_in):
    """
    Implement switch-like behavior.
    """

    # Here's some psuedocode to start you off implementing a learning
    # switch. You'll need to rewrite it as real Python code.

    # Learn the port for the source MAC
    #self.mac_to_port ... <add or update entry>
    self.mac_to_port.update({ packet.src: packet_in.in_port })

    log.debug( self.mac_to_port )

    if self.mac_to_port.get(packet.dst) is not None:

        # Send packet out the associated port
        self.resend_packet(packet_in, self.mac_to_port.get(packet.dst
        ))

        # Once you have the above working, try pushing a flow entry
        # instead of resending the packet (comment out the above and
        # uncomment and complete the below.)

        log.debug("Installing_flow...")
        # Maybe the log statement should have source/destination/port
        ?
        log.debug("Source:_" + str(packet.src) + "_Destination:_" +
            str(packet.dst) + "_Port:_" + str(self.mac_to_port.get(
            packet.dst)))

        msg = of.ofp_flow_mod()
        #msg.match = of.ofp_match.from_packet(packet)
        msg.match.in_port = packet_in.in_port
        msg.match.dl_dst = packet.dst
        msg.actions.append(of.ofp_action_output(port = self.

```

```

        mac_to_port.get(packet.dst)))

    self.connection.send(msg)
    #
    ## Set fields to match received packet
    #msg.match = of.ofp_match.from_packet(packet)
    #
    #< Set other fields of flow_mod (timeouts? buffer_id?) >
    #
    #< Add an output action, and send — similar to resend_packet
    () >

else:
    # Flood the packet out everything but the input port
    # This part looks familiar, right?
    self.resend_packet(packet_in, of.OFPP_ALL)

def _handle_PacketIn (self, event):
    """
    Handles packet in messages from the switch.
    """

    packet = event.parsed # This is the parsed packet data.
    if not packet.parsed:
        log.warning("Ignoring_incomplete_packet")
        return

    packet_in = event.ofp # The actual ofp_packet_in message.

    # Comment out the following line and uncomment the one after
    # when starting the exercise.
    #self.act_like_hub(packet, packet_in)
    self.act_like_switch(packet, packet_in)

def launch ():
    """
    Starts the component
    """
    def start_switch (event):
        log.debug("Controlling_%s" % (event.connection,))
        Tutorial(event.connection)
    core.openflow.addListenerByName("ConnectionUp", start_switch)

```